

E-mail: marcoantonio.stranisci@unito.it
ORCID: <https://orcid.org/0000-0001-9337-7250>
Google Scholar: [Marco Antonio Stranisci](#)
Personal website: <https://marcostranisci.github.io>

MARCO ANTONIO STRANISCI

Post-doc researcher at the University of Turin and co-founder of [aequa-tech](#).

APPLYING FOR: Fixed term Assistant Professor in the Department of “Computing Sciences” of “Università Commerciale Luigi Bocconi” of Milan.

PROFESSIONAL EXPERIENCE

Postdoctoral researcher - Università degli Studi di Torino

01/06/2024 - ongoing

I am involved in two research projects exploring how Knowledge Graphs (KGs) and Large Language Models (LLMs) can be integrated to support evidence-based policy making and enhance access to Cultural Heritage for diverse audiences and stakeholders. I co-authored the funding proposals for both projects, which successfully secured a combined total of €430,000.

Research and Business development - aequa-tech

01/07/2023 - ongoing

I am a co-founder of aequa-tech, a NLP start-up based in Turin, for where I lead business development. I am responsible for the company business model working closely with clients, companies, and academic institutions together with which I co-design our products and services.

Grant writing and research project manager - ACMOS

01/10/2018 - 28/02/2022

I secured a total of €530,000 in funding for two projects: Contro l'odio, funded by the Italian Ministry of Labour and BePositive!, funded by Google. Developed in collaboration with the University of Turin, these projects focused on the design and implementation of methods for the automatic detection of hate speech.

EDUCATION

PhD in Computer Science. Università degli Studi di Torino

01/10/2020 - 28/06/2024

Supervisor: Prof. Rossana Damiano. Co-supervisor: Prof. Viviana Patti

During my PhD, I combined Natural Language Processing (NLP) methods with Semantic Web technologies to design a framework for the detection and mitigation of biases in digital archives. My research has been accepted and awarded to the top-ranked conferences of both scientific

communities and one of the outcomes of my research received a grant of €150,000 that funded my Postdoctoral fellowship.

Visiting Student. IT University of Copenhagen

05/2024

Supervisor: Prof. Christian Hardmeier

My visit to the IT University of Copenhagen has been focused on preparing a proposal for the Marie Skłodowska-Curie (MSCA) Post-Doctoral Fellowship, which has received the seal of excellence in 2025. A survey paper correlated to the topic of the MSCA will be presented at the next AAAI conference (January, 2026, Singapore).

Visiting Student. University of Groningen

05/2022

Supervisor: Prog. Tommaso Caselli

During the research stay, I co-designed with Prof. Tommaso Caselli a research work that adopts biographical event detection to detect *representational* bias in Wikipedia English pages. The paper has been presented at the 61st edition of ACL (July, 2023, Toronto), where it was awarded an Outstanding Paper Award.

Master in Cognitive Linguistics. Università degli Studi di Torino

01/01/2010 - 12/11/2012

My primary interests have been Socio-linguistics, Pragmatics, and Cognitive Linguistics. My thesis focused on the creation of a corpus of conceptual metaphors extracted from political discourses.

SCIENTIFIC PRODUCTION

My scientific production is highly interdisciplinary and has been accepted at top-ranked conferences of different sectors of Computer Science: Natural Language Processing (ACL, EMNLP), Semantic Web (ISWC), and Artificial Intelligence (AAAI). Below, there is a selection of my publications showcasing my interdisciplinary profile and my ability to either lead or supervise research initiatives.

- 1)** Stranisci M.A., Hardmeier C.. (2026). **What Are They Filtering Out? A Survey of Filtering Strategies for Harm Reduction in Pretraining Datasets.** *Proceedings of the AAAI Conference on Artificial Intelligence*. AAAI Press.

CORE: A*

This work presents the first systematic analysis of the impact of data filtering strategies in the erasure of vulnerable identities from pretraining datasets. The research is the outcome of a collaboration with Christian Hardmeier, associate professor at the IT University of Copenhagen.

- 2)** Urbinati A., Lai M., Frenda S. & Stranisci M.A. (2025). **Are you sure? Measuring models bias in content moderation through uncertainty.** *Findings of the Association for Computational Linguistics: EMNLP 2025*. Association for Computational Linguistics.

CORE: A*

This work adopts the Conformal Prediction framework to assess the presence of biases in Language Models when they are adopted for the moderation of Hate Speech. I co-designed and supervised this research, which is an initiative of my start-up *aqua-tech*.

3) Lo S.M., Araque O., Sharma R. & Stranisci M.A. (2025). **That is Unacceptable: the Moral Foundations of Canceling.** In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. Association for Computational Linguistics.

CORE: A*

This work explores the impact of annotators' moral values in the perception of canceling attitudes in social media, adopting the Moral Foundations Theory. I co-designed the research and participated in the initiative with the role of coordinator.

4) Troiano E., Labat S., Stranisci M.A., Damiano R., Patti V., & Klinger R. (2024). **Dealing with Controversy: An Emotion and Coping Strategy Corpus Based on Role Playing.** In *Findings of the Association for Computational Linguistics: EMNLP 2024*. Association for Computational Linguistics.

CORE: A*

In this work we present the first corpus of coping strategies generated through a role playing experimental setting. The aim of the research is to provide a resource that can be used to observe how people cope with situations that stress their identities (e.g., anti-LGBTQIA claims). I led the initiative in collaboration with Enrica Troiano (University of Stuttgart) and Sofie Labat (Ghent University).

5) Stranisci M. A., Bernasconi E., Patti V., Ferilli S., Ceriani M., & Damiano, R. (2023). **The World Literature Knowledge Graph.** In *International Semantic Web Conference* (pp. 435-452). Springer Nature Switzerland. Springer Nature.

CORE: A

This work analyzes and mitigates allocative bias in Wikidata through the design of a semantic model for the alignment of knowledge from different sources. I led this research, which is one of the outputs of my PhD and has been granted with a funding of 150,000 euros by the Next Generation Internet Search Initiative.

6) Stranisci M. A., Damiano R., Mensa E., Patti V., Radicioni D., & Caselli T. (2023). **WikiBio: a Semantic Resource for the Intersectional Analysis of Biographical Events.** In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)* (pp. 12370-12384). Association for Computational Linguistics.

CORE: A*

This work presents the task of biographical event detection and uses it to explore the presence of *representational* bias in Wikipedia English biographies. I led this research, which is one of the outputs of my PhD and obtained an *Outstanding Paper Award* at the 61st Annual Meeting of the Association for Computational Linguistics.

INTERNATIONAL COLLABORATIONS

Between 2021 and 2025, I collaborated with more than 30 universities and research institutions from 14 different countries, and 4 continents. My membership in the "Knowledge Graphs in the Era of Large Language Models (KGELL)" COST Action and engagement with several EU initiatives (NGI, Euraxes), NGOs, and companies give multidimensionality to my international profile.

In recent months, I have started collaborations with additional universities, with results expected in the first half of 2026: University of Bonn (Prof. Lucie Flek) on human-value alignment; University of Hamburg (Prof. Anne Lauscher) on data filtering strategies; University of Gottingen (Prof. Lisa Beinborn), Ghent University (Dr. Pranaydeep Singh), and University of the Basque Country (Prof.

Rodrigo Agerri) on BabyLM; Fondazione Bruno Kessler (Dr. Marco Guerini) on argumentation and moral values; Università Milano Bicocca (Prof. Elisabetta Fersini) on the computational analysis of canceling; University of Gothenburg (Dr. Pierluigi Cassotti) on the pragmatic meaning of terms related to psychological disorder.

Below, I report a selection of academic events that contributed to the development of my international profile.

[POLAR @ SemEval-2026 \(July 2026, San Diego, USA\)](#). POLAR is a SemEval Task that will be co-located with the 64th edition of ACL. I am currently leading the preparation for the Italian dataset as an extension of an existing benchmark for Hate Speech and Homotransphobia detection

[The First Workshop on Social Context \(May 2026, Palma de Mallorca, Spain\)](#). This workshop will contribute to the definition of a shared framework for the study and processing of social context in NLP. The event will see the participation of internationally recognized researchers such as Prof. Maarten Sap (Carnegie Mellon University) and Prof. Roman Klinger (University of Bamberg).

[The First Information Disorder Workshop \(May 2026, Palma de Mallorca, Spain\)](#). The workshop aims to foster the interdisciplinary dialogue on the computational analysis of Information Disorder. Conceived as a research outcome of my start-up, the workshop foresees the participation of Prof. Claire Wardle (Cornell University), who has been the first to define this framework.

[Biases we Research by Tutorial \(May 2026, Palma de Mallorca, Spain\)](#). This tutorial will showcase methods from social psychology, semantic web, and applied mathematics for the detection and mitigation of bias, and will involve the participation of members of the Data Labelers Association, the first trade union of crowdworkers for machine learning in Africa. For the first time, crowdworkers will be invited to hold a tutorial at an NLP conference.

[The 9th Evaluation campaign of NLP and speech tools for Italian \(February 2026, Bari, Italy\)](#). I am one of the organizers of EVALITA, the periodic evaluation campaign of NLP systems for Italian. This event will include the participation of members of 15 universities from all over the country.

[The 8th Workshop on Natural Language for Artificial Intelligence \(NL4AI\) \(November 2024, Bolzano\)](#). NL4AI is the workshop of NLP co-located with the annual meeting of AlxIA. The aim of this recurring workshop is to collect contributions at the intersection of NLP and AI. I was one of the organizers of the 8th edition.

RESEARCH GRANT WRITING

Between 2018 and 2025 I actively participated with a prominent role in writing grants that obtained funding for more than 1 million of euros from local, national, and international entities. I am currently involved in a network of partners to present a proposal for the upcoming Horizon call for applications “CL2-2026-01-HERITAGE-06: Safeguarding linguistic diversity in Europe” in which I will lead 1 Work Package.

[HARMONIA \(2024-2025\)](#)

Funding institution: FAIR foundation

Lead organization: Università degli Studi di Torino

Duration: 18 months

Grant size: 223.401 €

Role: I co-led the preparation of the proposal with my former PhD supervisor, Prof. Viviana Patti. I prepared the first draft of the project and jointly worked with her on the following iterations on the proposal.

World Literature Knowledge Graph (2023-2024).

Funding institution: NGI Search Consortium

Lead organization: Università degli Studi di Torino

Duration: 1 year

Grant size: 150.000 €

Role: I led the writing of this project, which is the continuation of my PhD research on allocative bias in literary archives.

Debunker-Assistant (2023-2024).

Funding institution: NGI Search Consortium

Lead organization: aequa-tech

Grant size: 150.000 €

Duration: 1 year

Role: I co-wrote the proposal with the other members of the start-up. The grant has been a crucial contribution to develop the first prototype of our tool for the detection of online fake news.

BePositive! (2020-2022).

Funding institution: Google Impact Challenge

Lead organization: ACMOS

Grant size: 250.000 €

Duration: 2 years

Role: I co-wrote the proposal with the grant writing office of the non-profit where I worked before starting my PhD.

Contro l'odio (2018-2020).

Funding institution: Italian Ministry of Labor (*Ministero del Lavoro e delle Politiche Sociali*)

Lead organization: ACMOS

Grant size: 285.500 €

Duration: 2 years

Role: I co-wrote the proposal with the grant writing office of the non-profit where I worked before starting my PhD.

GRANTS AND AWARDS

MSCA Post-Doctoral Fellowship - Seal of Excellence (2025)

In 2024 I submitted a Marie Skłodowska-Curie Post-Doctoral Fellowship that received the seal of excellence with a score of 90.6/100. A revised version of the proposal is now under review.

DAAD Alnet Fellowship (2025)

In May 2025 I was selected as DAAD Alnet Fellow. The fellowship allowed me to engage with the German academic ecosystem. This experience allowed me to establish three fruitful international

collaborations: with Prof. Lucie Fleck from the University of Bonn, Prof. Lisa Beinborn from the University of Gottingen, and Prof. Anne Lauscher from the University of Hamburg.

Outstanding Paper Award - 61st Annual Meeting of ACL (2023)

I received an Outstanding Paper Award at the 61st Annual Meeting of ACL for the paper “WikiBio: a Semantic Resource for the Intersectional Analysis of Biographical Events”.

Best of 2021 Lab - 13th International Conference of the CLEF Association (2022)

The paper “Analysing Moral Beliefs for Detecting Hate Speech Spreaders on Twitter”, which presents a methodology that leverages on moral values to detect hate speech spreaders, was selected as Best of 2021 Lab at the 13th CLEF Conference.

CONFERENCE AND WORKSHOP PRESENTATIONS

INVITED

- 08/09/2025 [**Oral Presentation**](#) on the limitations of Named Entity Recognition technologies for Cultural Heritage during the Lorentz workshop “Enriching Digital Heritage with LLMs and Linked Open Data”. Organizer: Lorentz Foundation.
- 07/08/2025 [**Invited speaker**](#) at the University of Bonn (NLP colloquium) with the presentation “Exploring bias, explaining hate: two critical studies on harm detection in Natural Language Processing”. Organizer: Bonn University.
- 23/05/2025 [**Invited speaker**](#) at the Workshop “Truth is in the Eyes of the Machines” organized by the Jantina Tammes School of Digital Society, Technology and AI. Organizer: University of Groningen.
- 19/11/2024 [**Invited Speaker**](#) at the International Conference “Linked Open Data and Literary Studies”. Organizer: Free University of Berlin.
- 10/10/2023 [**Invited Speaker**](#) at the International Workshop “Wikipedia, Wikidata and Wikibase: Usage Scenarios for Literary Studies”. Organizer: Free University of Berlin.
- 03/07/2023 [**Guest talk**](#) at the workshop “Ontologies for Narrative and Fiction (online workshop)” where I presented the World Literature Knowledge Graph. Organizer: Stuttgart Media University.
- 20/01/2023 [**Invited Speaker**](#) at the “Japanese Visual Media Graph Workshop”. Organizer: Stuttgart Media University. Organizer: University of Groningen.
- 12/10/2019 [**Oral Presentation**](#) of the Contro l'odio project at the event “10 anni di talenti per la società civile”. Organizer: Fondazione Goria.

PEER-REVIEWED

- 01/2026 [**Poster presentation**](#) at the AAAI Conference on Artificial Intelligence: “What Are They Filtering Out? A Survey of Filtering Strategies for Harm Reduction in Pretraining Datasets”.
- 29/07/2025 [**Poster presentation**](#) at the 63st Annual Meeting of the Association of Computational Linguistics (ACL 2024): “*That is Unacceptable: the Moral Foundations of Canceling*”.
- 30/11/2023 [**Oral Presentation**](#) at the 9th Italian Conference on Computational Linguistics. Scientific

communication of the paper “*WikiBio: a Semantic Resource for the Intersectional Analysis of Biographical Events*”

08/11/2023 **Oral Presentation** at the 22nd International Semantic Web Conference (ISWC 2023): “*The World Literature Knowledge Graph*”.

09/07/2023 **Poster presentation** at the 61st Annual Meeting of the Association of Computational Linguistics (ACL 2023): “*WikiBio: a Semantic Resource for the Intersectional Analysis of Biographical Events*”

23/06/2022 **Oral presentation** at the 13th conference of Language Resources and Evaluation (LREC 2022): “*APPRReddit: a Corpus of Reddit Posts Annotated for Appraisal*”

20/06/2022 **Oral presentation** at the 18th ISA workshop co-located with the 13th conference of Language Resources and Evaluation: “Guidelines and a Corpus for Extracting Biographical Events”.

24/10/2021 **Oral presentation** at the 12th Workshop on Ontology Design Patterns: “Mapping Biographical Events to ODPs Through Lexico-Semantic Patterns”.

03/09/2021 **Oral presentation** at the 3rd conference on Language, Data, and Knowledge (LDK 2021): “*Representing the Under-Represented: a Dataset of Post-Colonial, and Migrant Writers*”

TEACHING AND TUTORING

TEACHING EXPERIENCE

1) Biases we Research by Tutorial (May, 2026). I ideated this tutorial on the detection and mitigation of bias through methods from social psychology, semantic web, and applied mathematics. The tutorial has been accepted to be co-located with the LREC conference in the upcoming year (May, 2026).

2) Guest lecturer for the MS Course on “Semantic Modeling” (November, 2025). I co-organized a 6-hour laboratory on Information Extraction and Linking for Cultural Heritage in the context of the course of Semantic Modeling held by professor Rossana Damiano. University of Turin - Department of Computer Science.

3) Tutor for the Master “Digital Communication Design” (January-July, 2025). I tutored the 2025 class of the master in “Digital Communication Design”. My responsibilities involved teaching Python, Javascript, SQL, PHP, HTML, and CSS. COREP

4) Guest lecturer for the Jean Monnet chair Com4T.EU (April-May, 2025). I organized a 8-hours laboratory on the use of Artificial Intelligence for tackling online disinformation. University of Turin - Department of Cultures, Politics and Society.

5) Guest lecturer for the MS Course on “Linguistic Resources for Natural Language Processing” (December, 2024). I curated an 8-hour annotation workshop on canceling and moral values. University of Turin - Department of Humanities and Department of Computer Science

TUTORING, MENTORING AND SUPERVISION

The tutoring, mentoring and supervision of the following students always resulted in a publication in a conference, workshop, or journal.

- **Arianna Longo, MSC student (2024). University of Turin.** Co-advisor of the thesis “Recognizing Vulnerable Identities in Italian and Spanish News Headlines: a Novel Corpus for a Fine-Grained Analysis of Hate Speech.”
- **Alessandro Saracco, MSC Student (2023). University of Turin.** Co-advisor of the thesis “Human-Robot Interaction by Ontology-Based Dialogue Engine”.
- **Ousmane Diakite, MSC Student (2023). University of Turin.** Co-advisor of the thesis “Estrazione di Eventi Biografici in Testo in Linguaggio naturale: il Caso degli Scrittori Sottorappresentati.”
- **Soda Marem Lo, MSC Student (2022). University of Pavia.** Tutor of the student’s external traineeship on the computational analysis of racial stereotypes in Twitter.
- **Michele De Leonardi, MSC Student (2022). University of Turin.** Co-advisor of the MSC thesis “Moral Foundations Theory e Social Media: Annotazione e Analisi di un Corpus di Dati Estratti da Twitter”.

ENTREPRENEURSHIP

Before and during my academic journey I had extensive experience in the intersection between research, entrepreneurship and social innovation. This led me to gain awareness on the impact that research can have in different social contexts and in the challenges related to engaging with different audiences.

Research for social innovation. My first contact with NLP and AI of technologies has been during my experience as member of the steering committee of ACMOS, an Italian non-profit organization for human rights. During this period I ideated the *Contro l’Odio* project in collaboration with the Computer Science Department of the University of Turin. *Contro l’odio* has been an initiative in which NLP-based technologies supported educational actions aimed at students and professionals. I played a central role in obtaining funding for more than €500,000 to kick-off and sustain the project between 2018 and 2022.

Business Modelling. By being admitted to the initiatives promoted by Google Impact Challenge, 2i3t Incubator, and NGI Search consortium, I developed crucial skills in business modelling with a specific focus on theory of change, business model patterns, and business plan design.

Business Development. Between 2023 and 2025 I raised more than €250,000 euros in grants and contracts for *aqua-tech*, the academic spinoff that I funded in 2023 together with other 5 researchers of the University of Turin. €90,000 of contracts from Università San Raffaele and Liguria Digitale will sustain the first half of 2026.

AI technologies for non-expert audiences. During the last year I have been siding with my scientific coordinator, Rossana Damiano, in the creation of a research laboratory on AI for Cultural Heritage

(CH) that enabled me to engage with different cultural institutions such as Museo Egizio, Museo Nazionale dell'Automobile, British Library, and Europeana. During this experience I understood the challenges related to the engagement with non-expert audiences in the co-design and evaluation of AI technologies.

Dissemination. I organized and participated in a high number of dissemination events on real-world applications of NLP technologies. In 2019 I organized the “Contro l’odio” workshop in which renowned figures engaged against discrimination have participated. E.g., Maria Laura Boldrini, Cathy La Torre, Federico Faloppa. Between 2020 and 2025 I have been invited as an expert speaker at conferences on hate speech, disinformation, and bias detection. I presented my experience as an entrepreneur on several occasions, including a podcast series organized by Euraxess and an event organized by the PhD Career Support initiative of the University of Turin.